


A Verified Implementation of B^+ -trees in



B⁺-trees


- Basis of modern database and file systems
- Ubiquitous and safety critical
- Non-trivial analysis

- balanced, sorted




Verified Implementation of B⁺-trees

- Verification of
 - isin/insert/delete
 - iterators
 - range queries
- Obtain executable imperative implementation



Refinement Approach

Functional Definition




- abstract reasoning
- no pointers

Refinement



Imperative Refinement



Refinement Approach

Functional Definition

```
datatype 'a tree =  
  Leaf ('a list) |  
  Node (('a tree × 'a) list)  
    ('a tree)
```



Imperative Refinement

```
datatype 'a treei =  
  Leafi ('a array) ('a treei ref option) |  
  Nodei (('a treei ref option × 'a) array)  
    ('a treei ref)
```

Refinement Approach

Functional Definition

```
datatype 'a tree =
  Leaf ('a list) |
  Node (('a tree × 'a) list)
    ('a tree)
```



Imperative Refinement

```
datatype 'a treei =
  Leafi ('a array) ('a treei ref option) |
  Nodei (('a treei ref option × 'a) array)
    ('a treei ref)
```

Refinement relation: $t \sim_t t_i : 'a \text{ tree} \Rightarrow 'a \text{ tree}_i \Rightarrow \text{bool}$

- $(\text{Node } xs) \sim_t (\text{Node}_i \text{ xs}_i) =$
 $\exists xs_{\text{int}}. xs_{\text{int}} \sim_A xs_i * \forall ((k, x), (k_i, x_i)) \in \text{zip}(xs, xs_{\text{int}}): k = k_i \wedge x \sim_t x_i$
- $(\text{Leaf } ks) \sim_t (\text{Leaf}_i \text{ ks}_i \text{ } -) = ks \sim_A ks_i$
- $\sim_t - = \text{false}$

Functional definition


```

fun isin: 'a tree  $\Rightarrow$  'a  $\Rightarrow$  bool where
  isin (Leaf ks) x = isin1 x ks |
  isin (Node xs t) x = case split xs x of
    (_, (sub, sep)#rs)  $\Rightarrow$  isin sub x |
    (_, [])  $\Rightarrow$  isin t x

```

assume: sorted(t)

show: $x \in \text{set}(t) \leftrightarrow \text{isin}(t, x)$



Refinement Approach

Functional Definition

```

fun isin: 'a tree  $\Rightarrow$  'a  $\Rightarrow$  bool where
  isin (Leaf ks) x = isin1 x ks |
  isin (Node xs t) x =
    case split xs x of
      (_, (sub, sep)#rs)  $\Rightarrow$  isin sub x |
      (_, [])  $\Rightarrow$  isin t x

```



Imperative Refinement

```

partial_fun (heap) isini: 'a treei ref  $\Rightarrow$  'a
 $\Rightarrow$  bool Heap where
  isini p x = do {
    node  $\leftarrow$  !p;
    case node of {
      Leafi ks _  $\Rightarrow$  isin1i ks x |
      Nodei xs t  $\Rightarrow$  do {
        i  $\leftarrow$  spliti xs t;
        xsl  $\leftarrow$  length xs;
        if i < xsl then do {
          s  $\leftarrow$  xs[i];
          let (sub, sep) = s in
            isini (the sub) x
        } else isini t x
      }
    }
  }

```


Refinement approach

assume: $\text{sorted}(t), \text{order}_k(t)$

show:

$\langle t \sim_t t_i \rangle$

$\text{isin}_i t_i x$

$\langle \lambda \text{ res. } t \sim_t t_i * \text{isin}(t, x) \sim_b \text{res} \rangle$

Refinement approach

assume: $\text{sorted}(t), \text{order}_k(t)$

show:

$\langle t \sim_t t_i \rangle$

$\text{isin}_i t_i x$

$\langle \lambda \text{ res}. t \sim_t t_i * \text{isin}(t, x) \sim_b \text{res} \rangle$

show:

$\langle t \sim_t t_i \rangle$

$\text{insert}_i t_i x$

$\langle \lambda \text{ res}. \text{insert}(t, x) \sim_t \text{res} \rangle$

Refinement relation

Recall $t \sim t_i$:

- $(\text{Node } xs) \sim_t (\text{Node}_i xs_i) =$
 $\exists xs_{int}. xs_{int} \sim_A xs_i * \forall ((k, x), (k_i, x_i)) \in \text{zip}(xs, xs_{int}): k = k_i \wedge x \sim_t x_i$
- $(\text{Leaf } ks) \sim_t (\text{Leaf}_i ks_i _) = ks \sim_A ks_i$
- $_ \sim_t _ = \text{false}$

Recursive “view” from the root

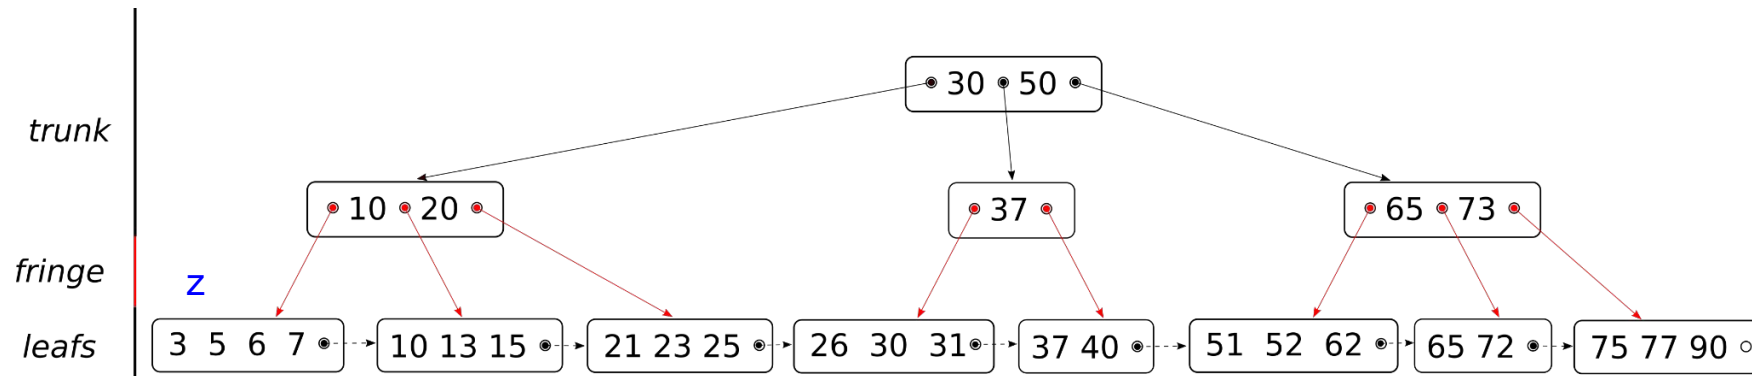
→ Need alternative when expressing iterators

Iterators

Ideally: $t \sim_t t_i \leftrightarrow \text{leafs}(t) \sim_L \mathbf{z} * \text{trunk}(t) \sim_T t_i$

Obstacle:

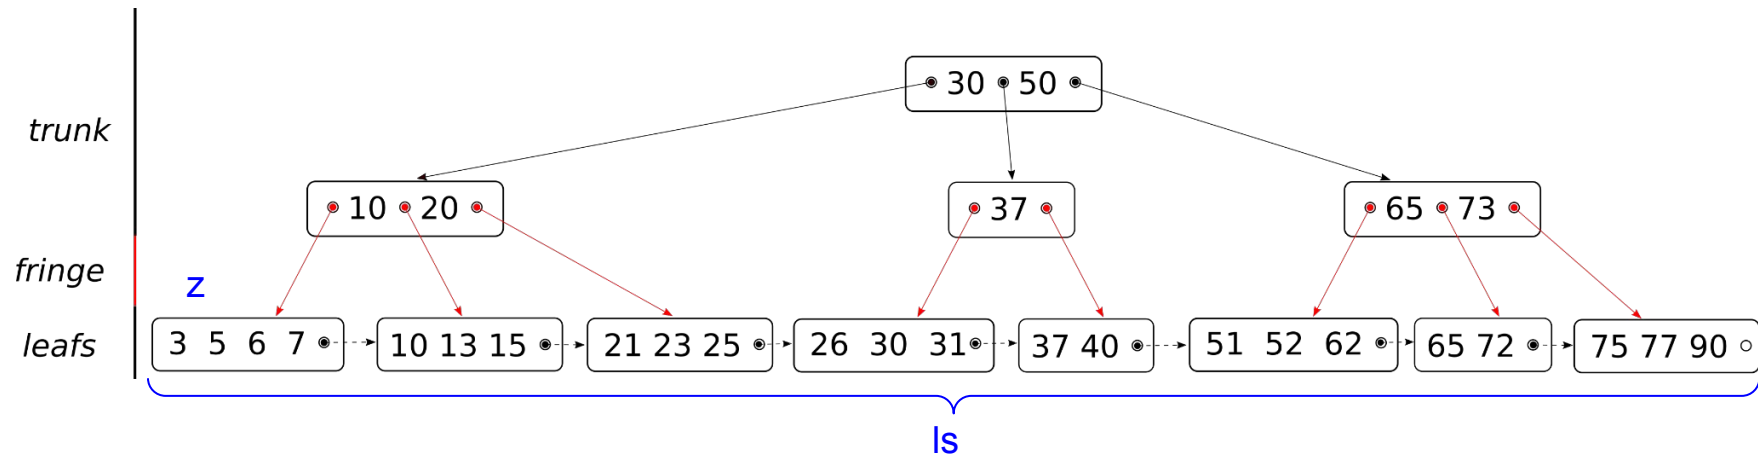
- need to show that tree can be split and re-combined



Iterators

Solution: $t \sim_t t_i \leftrightarrow \exists \text{ls}. \text{leafs}(t) \sim_L(z, \text{ls}) * \text{trunk}(t) \sim_T(t_i, \text{ls})$

Ghost variable **ls** pins the fringe pointers



Iterators

Implement standard iterator interface from list equivalence


$$\text{state}(x, z, t) = \text{trunk}(t) \sim_{\text{T}} t_i * \underbrace{\exists xs. \text{leafs}(t) = zs @ xs * zs \sim_{\text{L}} z * xs \sim_{\text{L}}}_{\text{current position of the iterator}}$$

- 1) $\langle t \sim_{\text{T}} t_i \rangle \text{init}(t) \langle \lambda z. \text{state}(z, z, t) \rangle$
- 2) $\text{has_next}(x): 'a \text{ tree}_i \Rightarrow \text{bool}$
- 3) $\text{next}(x): 'a \text{ tree}_i \Rightarrow ('a, 'a \text{ tree}_i)$
- 4) $\text{state}(_, z, t) \Rightarrow t \sim_{\text{T}} t_i$

Range queries

$$\text{lrange}(t, x) = \{ y \mid y \in \text{set}(t) \wedge y \geq x \}$$


Idea: Efficiently obtain $\text{state}(x, z, t)$ s.t. $\text{next}(x)$ successively returns $\text{lrange}(t, x)$



Range queries

Refinement steps:


- 1) straightforward functional *lrange* implementation



Range queries

Refinement steps:

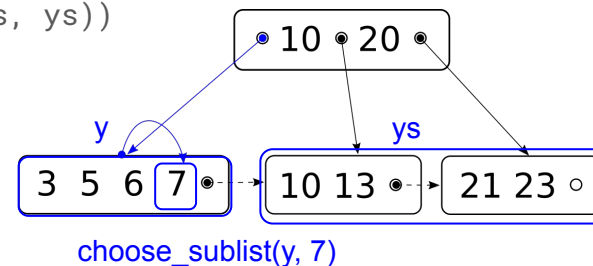
1) straightforward functional *lrange* implementation



2) *lrange_list*: collect leaf nodes and flatten later

a) `lrange_list(t, x) = y#ys : [Leaf]`

b) `lrange(t, x) = choose_sublist(y, x) @ concat(map(elements, ys))`



Range queries

Refinement steps:

1) straightforward functional *lrange* implementation

2) *lrange_list*: collect leaf nodes and flatten later



a) $\text{lrange_list}(t, x) = y\#ys : [\text{Leaf}]$

b) $\text{lrange}(t, x) = \text{choose_sublist}(y, x) @ \text{concat}(\text{map}(\text{elements}, ys))$

3) refine *lrange_list* imperatively

a) $\langle t \sim_t t_i \rangle$
 $\text{lrange_list}_i(t_i, x)$
 $\langle \lambda y_i. \text{trunk}(t) \sim_T t_i * \exists xs. \text{leafs}(t) = xs @ \text{lrange_list}(t, x) * xs \sim_L z * \text{lrange_list}(t, x) \sim_L y_i \rangle$

b) transform post condition to desired iterator state (see iterator state def.)



Related work

- Pen and paper by Fielding [TR 1980] and Sexton and Thielecke [MFPS 2008]
 - gradual refinement
 - separation logic
- Ernst *et al.* (KIV/TVLA) [SSM 2015]
 - direct imperative verification
 - improved shape analysis, without linked leafs
- Malecha *et al.* (Coq) [POPL 2010]
 - tree including linked leafs, iterator
 - refinement based

Evaluation

- strict separation of functional and imperative layers
- modular design, allowing for substitution with more efficient implementations

	Malecha <i>et.al.</i> ⁺	Ernst <i>et.al.</i> ^d	Our approach ⁺
Functional	360	-	413
Imperative	510	1862	1093
Proofs	5190	350 + 510 + 2940	6432 + 2231
Time (months)		6+	6 + 6

⁺ implementation of iterators, ^d implementation of deletion

Conclusion

- formal verification of imperative B⁺-trees succeeded
- implemented and verified first efficient range query
- including binary search for node internal navigation

